

In [123]:

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import warnings
warnings.filterwarnings("ignore")
from sklearn.model_selection import train_test_split, RandomizedSearchCV
from sklearn.linear_model import LinearRegression, LogisticRegression, RidgeClassifier
from sklearn.neighbors import KNeighborsRegressor
from sklearn.metrics import confusion_matrix, classification_report, accuracy_score
from sklearn.ensemble import RandomForestClassifier
```

In [2]:

```
df = pd.read_csv("water_potability.csv")
df.head()
```

Out[2]:

	ph	Hardness	Solids	Chloramines	Sulfate	Conductivity	Organic_carbon
0	NaN	204.890455	20791.318981	7.300212	368.516441	564.308654	10.37971
1	3.716080	129.422921	18630.057858	6.635246	NaN	592.885359	15.1800
2	8.099124	224.236259	19909.541732	9.275884	NaN	418.606213	16.8686
3	8.316766	214.373394	22018.417441	8.059332	356.886136	363.266516	18.4365
4	9.092223	181.101509	17978.986339	6.546600	310.135738	398.410813	11.5582

EDA

In [3]:

```
df.shape
```

Out[3]:

(3276, 10)

In [4]:

```
df.isna().sum()
```

Out[4]:

```
ph                491
Hardness          0
Solids            0
Chloramines       0
Sulfate           781
Conductivity      0
Organic_carbon    0
Trihalomethanes   162
Turbidity         0
Potability        0
dtype: int64
```

In [5]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3276 entries, 0 to 3275
Data columns (total 10 columns):
#   Column                Non-Null Count  Dtype  
---  -
0   ph                    2785 non-null   float64
1   Hardness              3276 non-null   float64
2   Solids                3276 non-null   float64
3   Chloramines           3276 non-null   float64
4   Sulfate               2495 non-null   float64
5   Conductivity          3276 non-null   float64
6   Organic_carbon        3276 non-null   float64
7   Trihalomethanes       3114 non-null   float64
8   Turbidity             3276 non-null   float64
9   Potability            3276 non-null   int64   
dtypes: float64(9), int64(1)
memory usage: 256.1 KB
```

In [6]:

```
df.columns
```

Out[6]:

```
Index(['ph', 'Hardness', 'Solids', 'Chloramines', 'Sulfate', 'Conductivity',
      'Organic_carbon', 'Trihalomethanes', 'Turbidity', 'Potability'],
      dtype='object')
```

working on null values

In [7]:

```
df.isna().sum()
```

Out[7]:

```
ph                491
Hardness           0
Solids             0
Chloramines        0
Sulfate           781
Conductivity       0
Organic_carbon     0
Trihalomethanes    162
Turbidity          0
Potability         0
dtype: int64
```

In [8]:

```
df["ph"].fillna(df["ph"].mean(),inplace=True)
```

In [9]:

```
df["Sulfate"].fillna(df["Sulfate"].mean(), inplace=True)
```

In [10]:

```
df["Trihalomethanes"].fillna(df["Trihalomethanes"].mean(),inplace=True)
```

In [11]:

```
df.isna().sum()
```

Out[11]:


```
ph                0
Hardness           0
Solids             0
Chloramines        0
Sulfate            0
Conductivity       0
Organic_carbon     0
Trihalomethanes    0
Turbidity          0
Potability         0
dtype: int64
```

In [12]:

```
df.describe()
```

Out[12]:

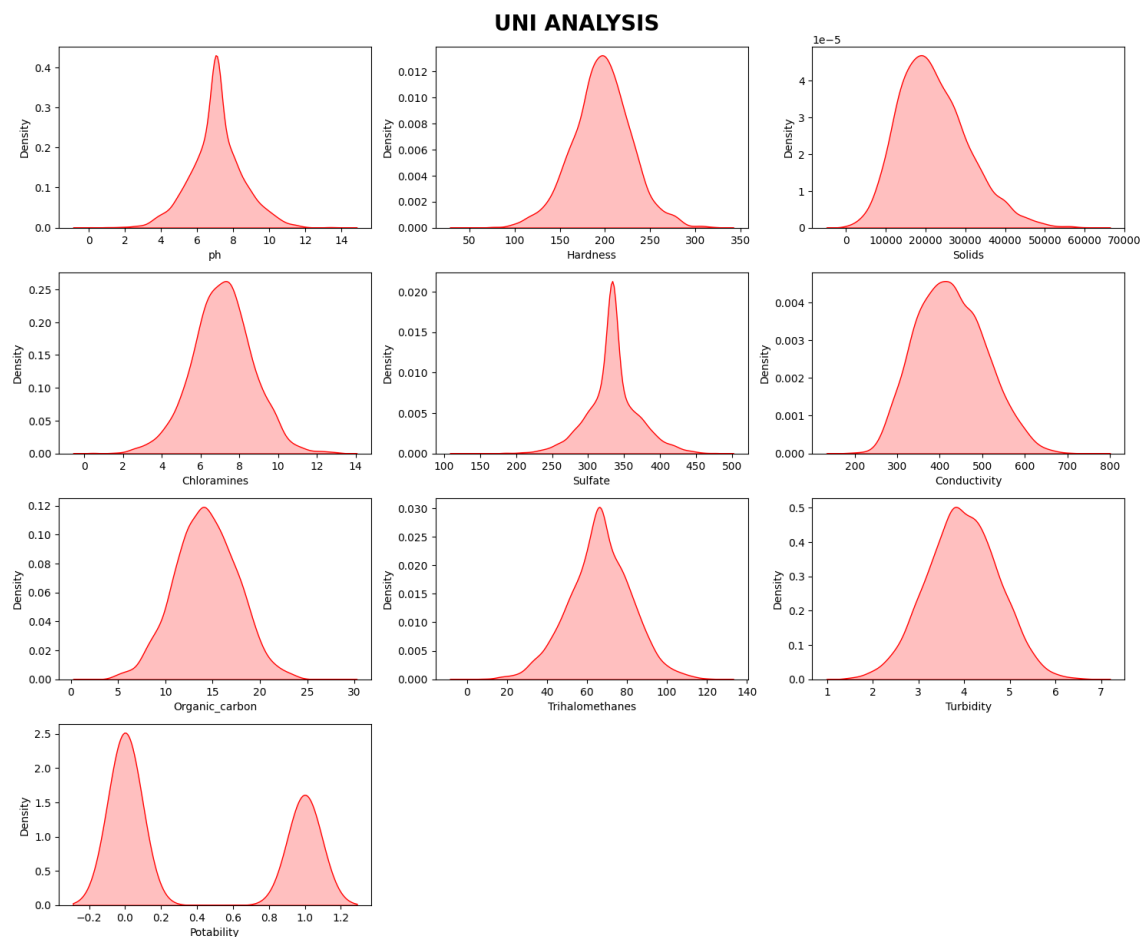
	ph	Hardness	Solids	Chloramines	Sulfate	Conductivity	Org
count	3276.000000	3276.000000	3276.000000	3276.000000	3276.000000	3276.000000	
mean	7.080795	196.369496	22014.092526	7.122277	333.775777	426.205111	
std	1.469956	32.879761	8768.570828	1.583085	36.142612	80.824064	
min	0.000000	47.432000	320.942611	0.352000	129.000000	181.483754	
25%	6.277673	176.850538	15666.690297	6.127421	317.094638	365.734414	
50%	7.080795	196.967627	20927.833607	7.130299	333.775777	421.884968	
75%	7.870050	216.667456	27332.762127	8.114887	350.385756	481.792304	
max	14.000000	323.124000	61227.196008	13.127000	481.030642	753.342620	



univariate analysis

In [13]:

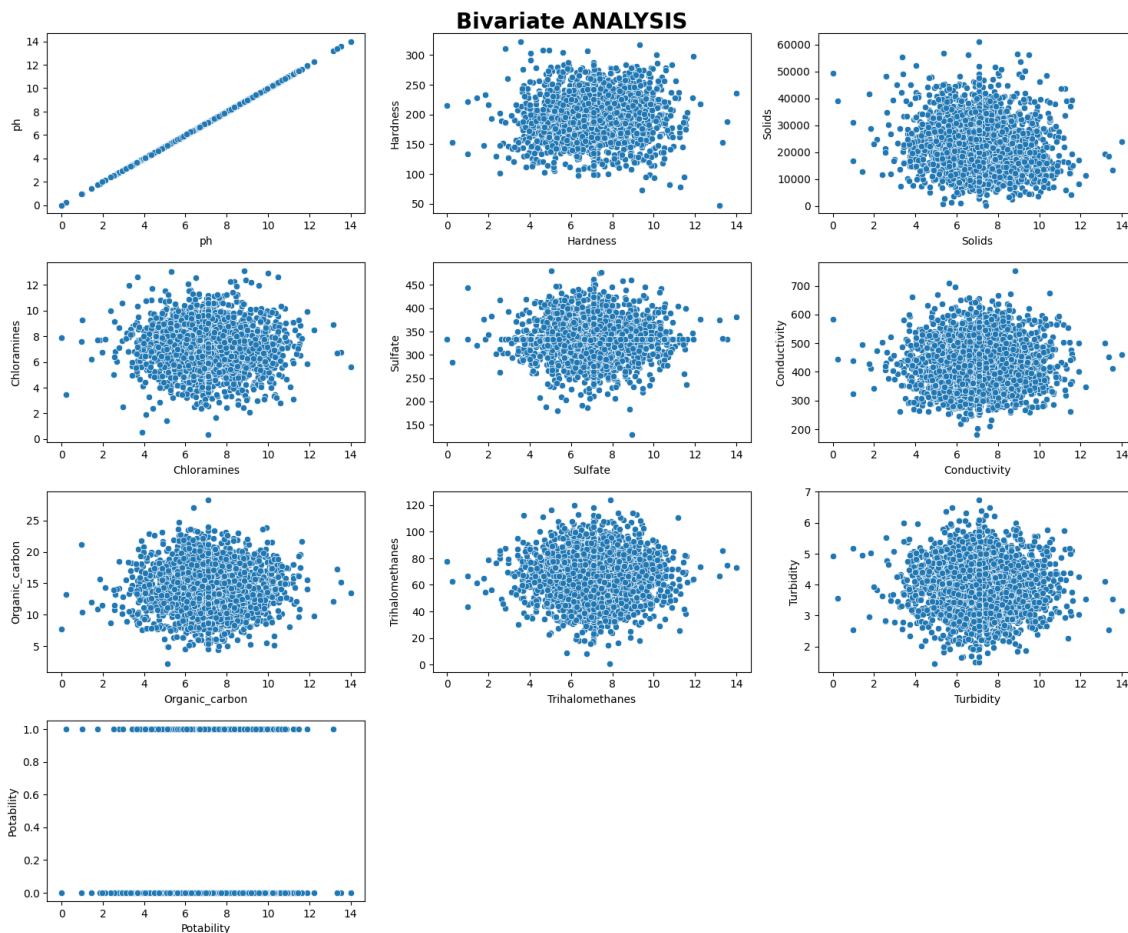
```
plt.figure(figsize=(15,15))
plt.suptitle("UNI ANALYSIS", fontsize=20 , fontweight="bold")
for i in range(len(df.columns)):
    plt.subplot(5,3,i+1)
    sns.kdeplot(x=df[df.columns[i]], shade=True , color = "r")
    plt.xlabel(df.columns[i])
plt.tight_layout()
```



Bivariate Analysis

In [14]:

```
plt.figure(figsize=(15,15))
plt.suptitle("Bivariate ANALYSIS", fontsize=20, fontweight="bold")
for i in range(len(df.columns)):
    plt.subplot(5,3,i+1)
    sns.scatterplot(x=df["ph"], y=df[df.columns[i]])
    plt.xlabel(df.columns[i])
    plt.tight_layout()
```



In [15]:

```
df.columns
```

Out[15]:

```
Index(['ph', 'Hardness', 'Solids', 'Chloramines', 'Sulfate', 'Conductivity',
      'Organic_carbon', 'Trihalomethanes', 'Turbidity', 'Potability'],
      dtype='object')
```

INSIGHTS


In [16]:

```
import plotly.express as px
```

In [17]:

```
fig = px.histogram(x= df["ph"], y=df["Hardness"], title= "Hardness contained in water w.  
fig.show()
```


Hardness contained in water w.rto ph



In [18]:

```
fig = px.histogram(x= df["ph"], y=df["Chloramines"], title= "Chloramines contained in wa  
fig.show()
```


Chloramines contained in water w.rto ph



In [19]:

```
fig = px.histogram(x= df["ph"], y=df["Conductivity"], title= "Conductivity contained in  
fig.show()
```

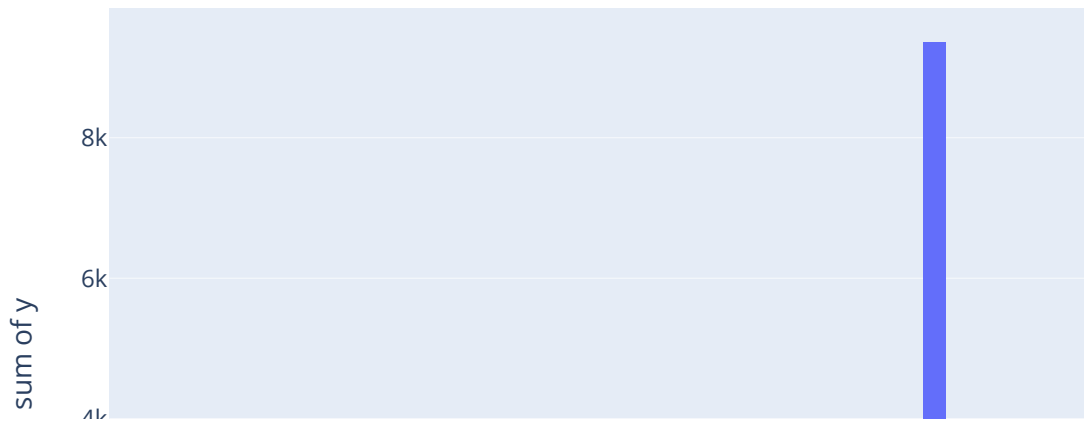
Conductivity contained in water w.rto ph



In [20]:

```
fig = px.histogram(x= df["ph"], y=df["Organic_carbon"], title= "Organic_carbon contained  
fig.show()
```

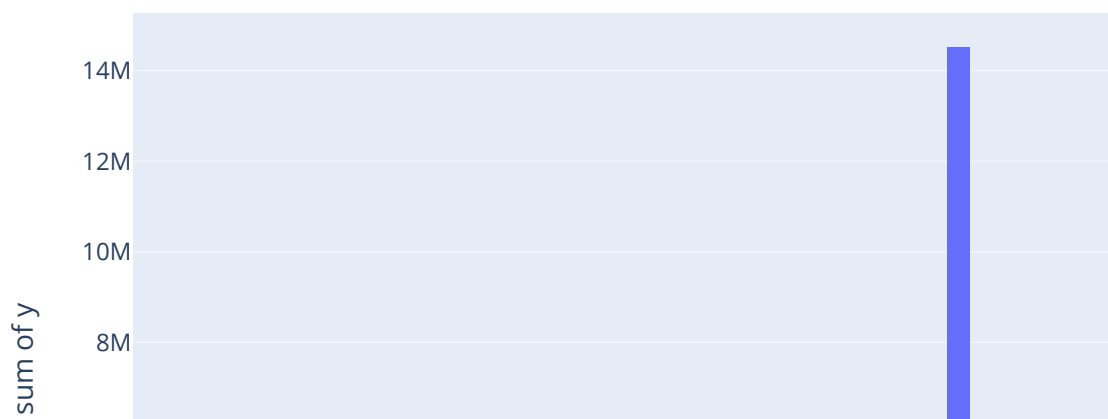
Organic_carbon contained in water w.rto ph



In [21]:

```
fig = px.histogram(x= df["ph"], y=df["Solids"], title= "Solids contained in water w.rto  
fig.show()
```

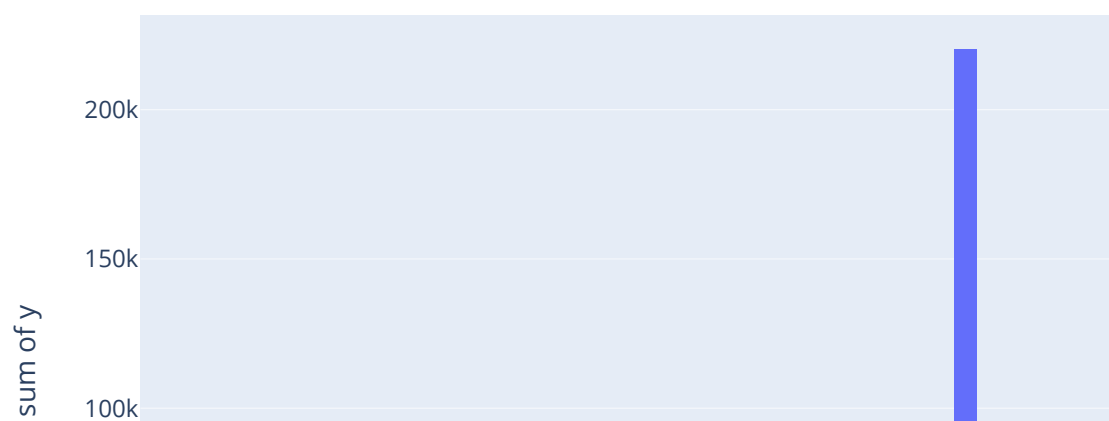
Solids contained in water w.rto ph



In [22]:

```
fig = px.histogram(x= df["ph"], y=df["Sulfate"], title= "Sulfate contained in water w.rto  
fig.show()
```

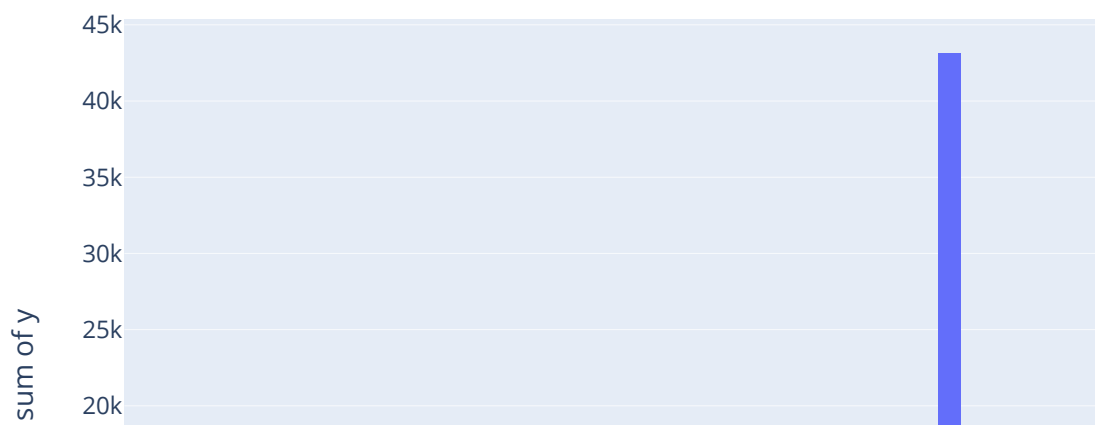
Sulfate contained in water w.rto ph



In [23]:

```
fig = px.histogram(x= df["ph"], y=df["Trihalomethanes"], title= "Trihalomethanes contain  
fig.show()
```


Trihalomethanes contained in water w.rto ph



In [24]:

```
fig = px.histogram(x= df["ph"], y=df["Turbidity"], title= "Turbidity contained in water",  
fig.show()
```

Turbidity contained in water w.rto ph



In [25]:

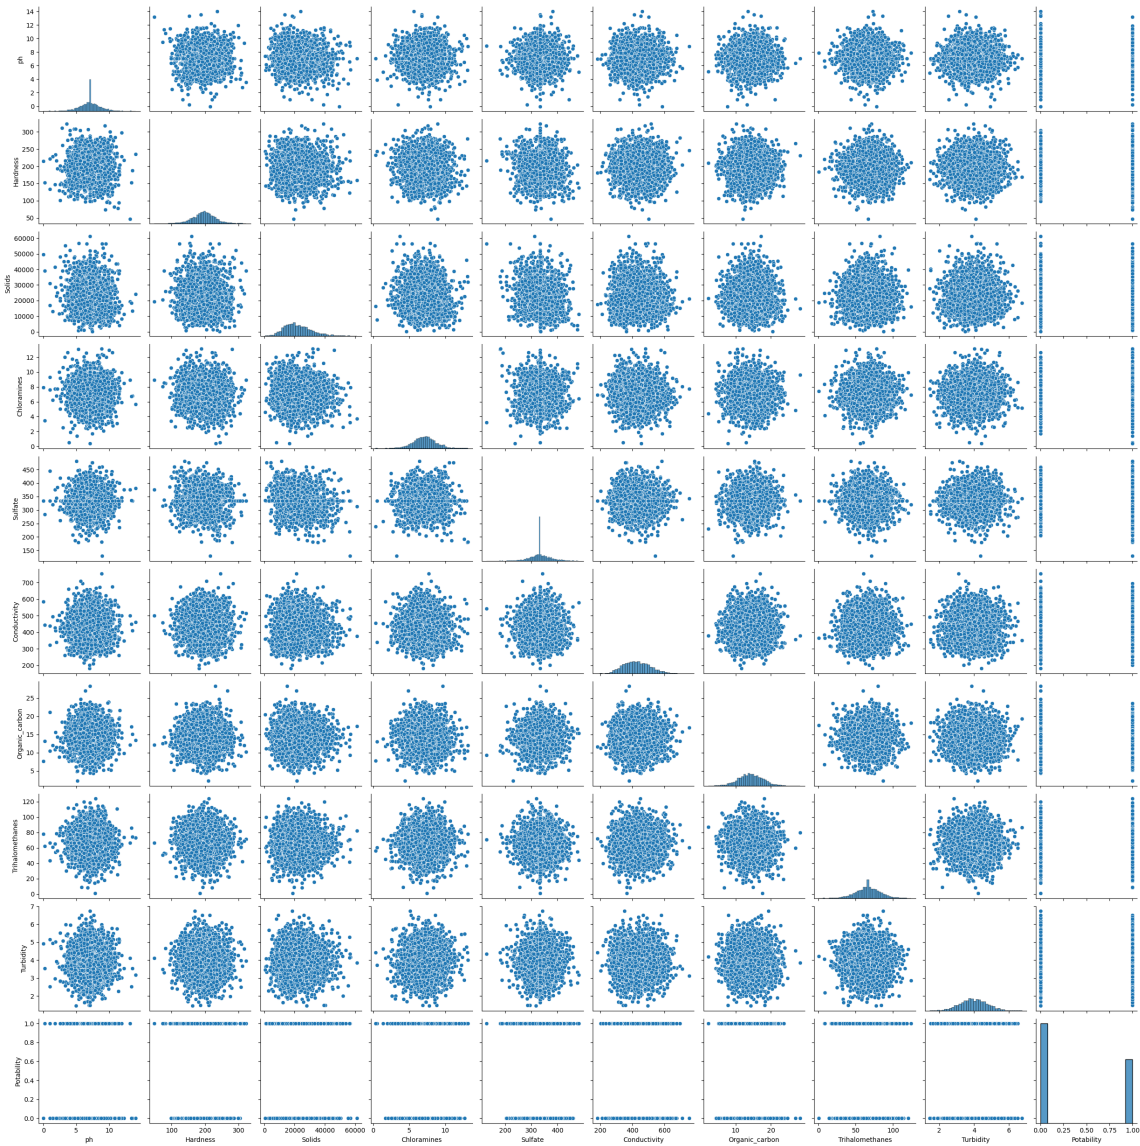
```
# fig = px.bar(df,df["Potability"] , y= df["ph"])  
# fig.show()
```

In [26]:

```
sns.pairplot(df)
```

Out[26]:

<seaborn.axisgrid.PairGrid at 0x1f75f3d7130>




In [27]:

```
sns.countplot(x= df["Potability"])
```

Out[27]:

<AxesSubplot: xlabel='Potability', ylabel='count'>



Target column is high

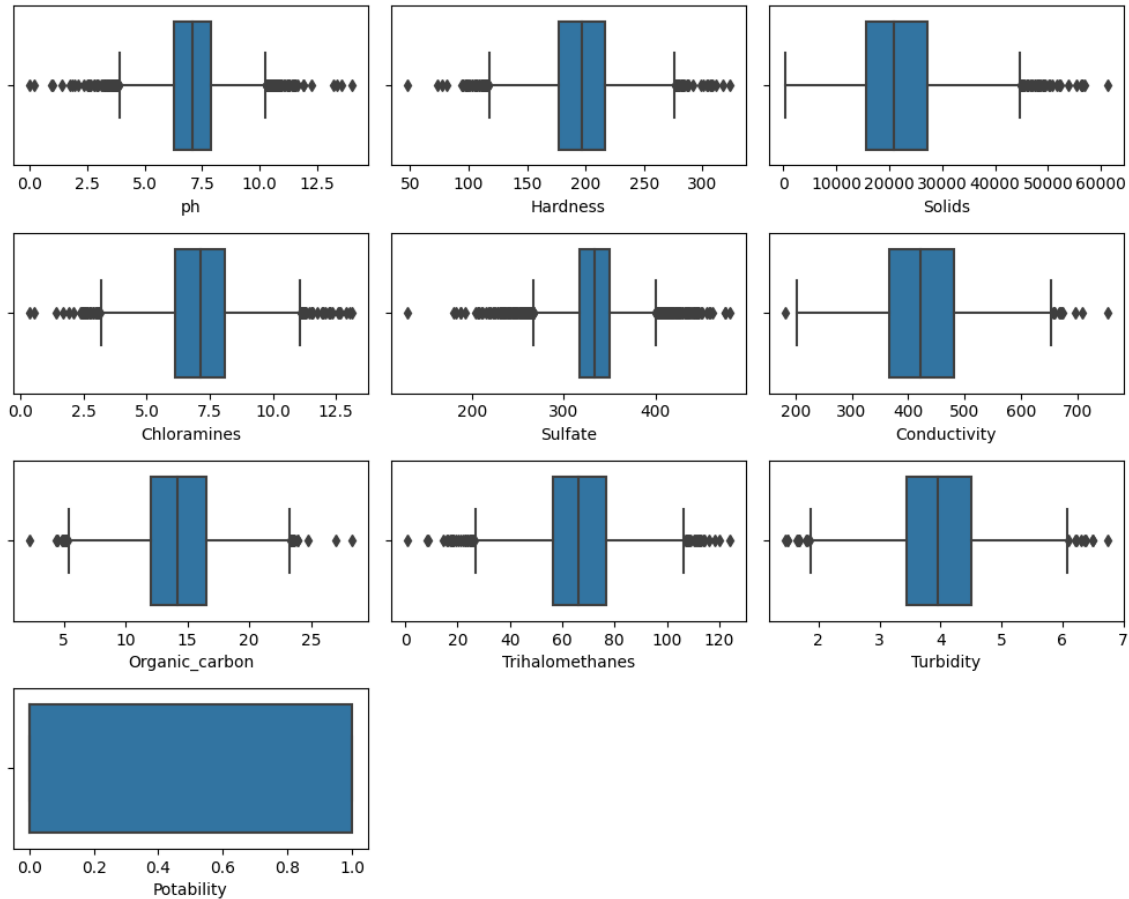
Feature Engineering

Detecting Outliers

In [28]:

```
plt.figure(figsize=(10,10))
plt.suptitle("Outliers Analysis", fontsize=20, fontweight="bold")
for i in range(0,len(df.columns)):
    plt.subplot(5,3,i+1)
    sns.boxplot(x=df[df.columns[i]])
    plt.xlabel(df.columns[i])
    plt.tight_layout()
```

Outliers Analysis




ph

In [29]:

```
sns.boxplot(df["ph"])
```

Out[29]:

<AxesSubplot: >



In [30]:

```
df["ph"].describe()
```

Out[30]:

```
count    3276.000000
mean       7.080795
std        1.469956
min         0.000000
25%        6.277673
50%        7.080795
75%        7.870050
max        14.000000
Name: ph, dtype: float64
```

In [31]:

```
print(df["ph"].quantile(0.10))
print(df["ph"].quantile(0.90))
```

```
5.282194491912188
8.925046875415749
```

In [32]:


```
df["ph"] = np.where(df["ph"] < 5.282194491912188, 5.282194491912188, df["ph"])  
df["ph"] = np.where(df["ph"] > 8.925046875415749, 8.925046875415749, df["ph"])
```

In [33]:

```
sns.boxplot(df["ph"])
```

Out[33]:

<AxesSubplot: >




Hardness

In [34]:

```
sns.boxplot(df["Hardness"])
```

Out[34]:

<AxesSubplot: >



In [35]:

```
df["Hardness"].describe()
```

Out[35]:

```
count    3276.000000
mean      196.369496
std        32.879761
min        47.432000
25%       176.850538
50%       196.967627
75%       216.667456
max       323.124000
Name: Hardness, dtype: float64
```

In [36]:

```
print(df["Hardness"].quantile(0.10))
print(df["Hardness"].quantile(0.90))
```

```
155.2239641077801
236.35070740017414
```

In [37]:


```
df["Hardness"] = np.where(df["Hardness"] < 155.2239641077801, 155.2239641077801, df["Hardness"])  
df["Hardness"] = np.where(df["Hardness"] > 236.35070740017414, 236.35070740017414, df["Hardness"])
```

In [38]:

```
sns.boxplot(df["Hardness"])
```

Out[38]:

<AxesSubplot: >




Solids

In [39]:

```
sns.boxplot(df["Solids"])
```

Out[39]:

<AxesSubplot: >



In [40]:

```
df["Solids"].describe()
```

Out[40]:

```
count      3276.000000
mean       22014.092526
std         8768.570828
min         320.942611
25%        15666.690297
50%        20927.833607
75%        27332.762127
max         61227.196008
Name: Solids, dtype: float64
```

In [41]:

```
print(df["Solids"].quantile(0.10))
print(df["Solids"].quantile(0.90))
```

```
11740.528189473214
33814.93523020222
```

In [42]:


```
df["Solids"] = np.where(df["Solids"] < 11740.528189473214, 11740.528189473214, df["Solids"])  
df["Solids"] = np.where(df["Solids"] > 33814.93523020222, 33814.93523020222, df["Solids"])
```

In [43]:

```
sns.boxplot(df["Solids"])
```

Out[43]:

<AxesSubplot: >



In [44]:

```
df.columns
```

Out[44]:

```
Index(['ph', 'Hardness', 'Solids', 'Chloramines', 'Sulfate', 'Conductivity',  
      'Organic_carbon', 'Trihalomethanes', 'Turbidity', 'Potability'],  
      dtype='object')
```

Chloramines

In [45]:

```
sns.boxplot(df["Chloramines"])


print(df["Chloramines"].quantile(0.10))
print(df["Chloramines"].quantile(0.90))
df["Chloramines"].describe()
```

5.181270677724534

9.122578323075329

Out[45]:

```
count    3276.000000
mean       7.122277
std        1.583085
min         0.352000
25%         6.127421
50%         7.130299
75%         8.114887
max        13.127000
Name: Chloramines, dtype: float64
```




In [46]:

```
df["Chloramines"] = np.where(df["Chloramines"] < 5.181270677724534, 5.181270677724534, df["Chloramines"])  
df["Chloramines"] = np.where(df["Chloramines"] > 9.122578323075329, 9.122578323075329, df["Chloramines"])  
sns.boxplot(df["Chloramines"])
```

Out[46]:

<AxesSubplot: >



Sulfate

In [47]:


```
sns.boxplot(df["Sulfate"])\n\nprint(df["Sulfate"].quantile(0.10))\nprint(df["Sulfate"].quantile(0.90))\ndf["Sulfate"].describe()
```

290.055010521933

378.4783210497187

Out[47]:

```
count    3276.000000\nmean      333.775777\nstd       36.142612\nmin       129.000000\n25%       317.094638\n50%       333.775777\n75%       350.385756\nmax       481.030642\nName: Sulfate, dtype: float64
```




In [48]:

```
df["Sulfate"] = np.where(df["Sulfate"] < 290.055010521933, 290.055010521933, df["Sulfate"])  
df["Sulfate"] = np.where(df["Sulfate"] > 378.4783210497187, 378.4783210497187, df["Sulfate"])  
sns.boxplot(df["Sulfate"])
```

Out[48]:

<AxesSubplot: >



Conductivity

In [49]:


```
sns.boxplot(df["Conductivity"])

print(df["Conductivity"].quantile(0.10))
print(df["Conductivity"].quantile(0.90))
df["Conductivity"].describe()
```

```
325.1171240676143
533.2972414189196
```

Out[49]:

```
count    3276.000000
mean      426.205111
std       80.824064
min       181.483754
25%       365.734414
50%       421.884968
75%       481.792304
max       753.342620
Name: Conductivity, dtype: float64
```




In [50]:

```
df["Conductivity"] = np.where(df["Conductivity"] < 325.1171240676143, 325.1171240676143, df["Conductivity"])  
df["Conductivity"] = np.where(df["Conductivity"] > 533.2972414189196, 533.2972414189196, df["Conductivity"])  
sns.boxplot(df["Conductivity"])
```

Out[50]:

<AxesSubplot: >



Organic_carbon

In [51]:

```
sns.boxplot(df["Organic_carbon"])


print(df["Organic_carbon"].quantile(0.10))
print(df["Organic_carbon"].quantile(0.90))
df["Organic_carbon"].describe()
```

10.123765383583503

18.50456708562831

Out[51]:

```
count    3276.000000
mean      14.284970
std        3.308162
min        2.200000
25%       12.065801
50%       14.218338
75%       16.557652
max       28.300000
Name: Organic_carbon, dtype: float64
```




In [52]:

```
df["Organic_carbon"] = np.where(df["Organic_carbon"] < 10.123765383583503, 10.1237653835  
df["Organic_carbon"] = np.where(df["Organic_carbon"] > 18.50456708562831, 18.5045670856283  
sns.boxplot(df["Organic_carbon"])
```

Out[52]:

<AxesSubplot: >



Trihalomethanes

In [53]:

```
sns.boxplot(df["Trihalomethanes"])


print(df["Trihalomethanes"].quantile(0.10))
print(df["Trihalomethanes"].quantile(0.90))
df["Trihalomethanes"].describe()
```

46.209472917430155

85.90009466879661

Out[53]:

```
count    3276.000000
mean       66.396293
std       15.769881
min        0.738000
25%       56.647656
50%       66.396293
75%       76.666609
max      124.000000
Name: Trihalomethanes, dtype: float64
```




In [54]:

```
df["Trihalomethanes"] = np.where(df["Trihalomethanes"] < 46.209472917430155, 46.20947291  
df["Trihalomethanes"] = np.where(df["Trihalomethanes"] > 85.90009466879661, 85.90009466879  
sns.boxplot(df["Trihalomethanes"])
```

Out[54]:

<AxesSubplot: >



Turbidity

In [55]:

```
sns.boxplot(df["Turbidity"])


print(df["Turbidity"].quantile(0.10))
print(df["Turbidity"].quantile(0.90))
df["Turbidity"].describe()
```

2.9518028252699757

4.977140682806077

Out[55]:

```
count    3276.000000
mean      3.966786
std       0.780382
min       1.450000
25%       3.439711
50%       3.955028
75%       4.500320
max       6.739000
Name: Turbidity, dtype: float64
```




In [56]:

```
df["Turbidity"] = np.where(df["Turbidity"] < 2.9518028252699757, 2.9518028252699757, df["Turbidity"])  
df["Turbidity"] = np.where(df["Turbidity"] > 4.977140682806077, 4.977140682806077, df["Turbidity"])  
sns.boxplot(df["Turbidity"])
```

Out[56]:

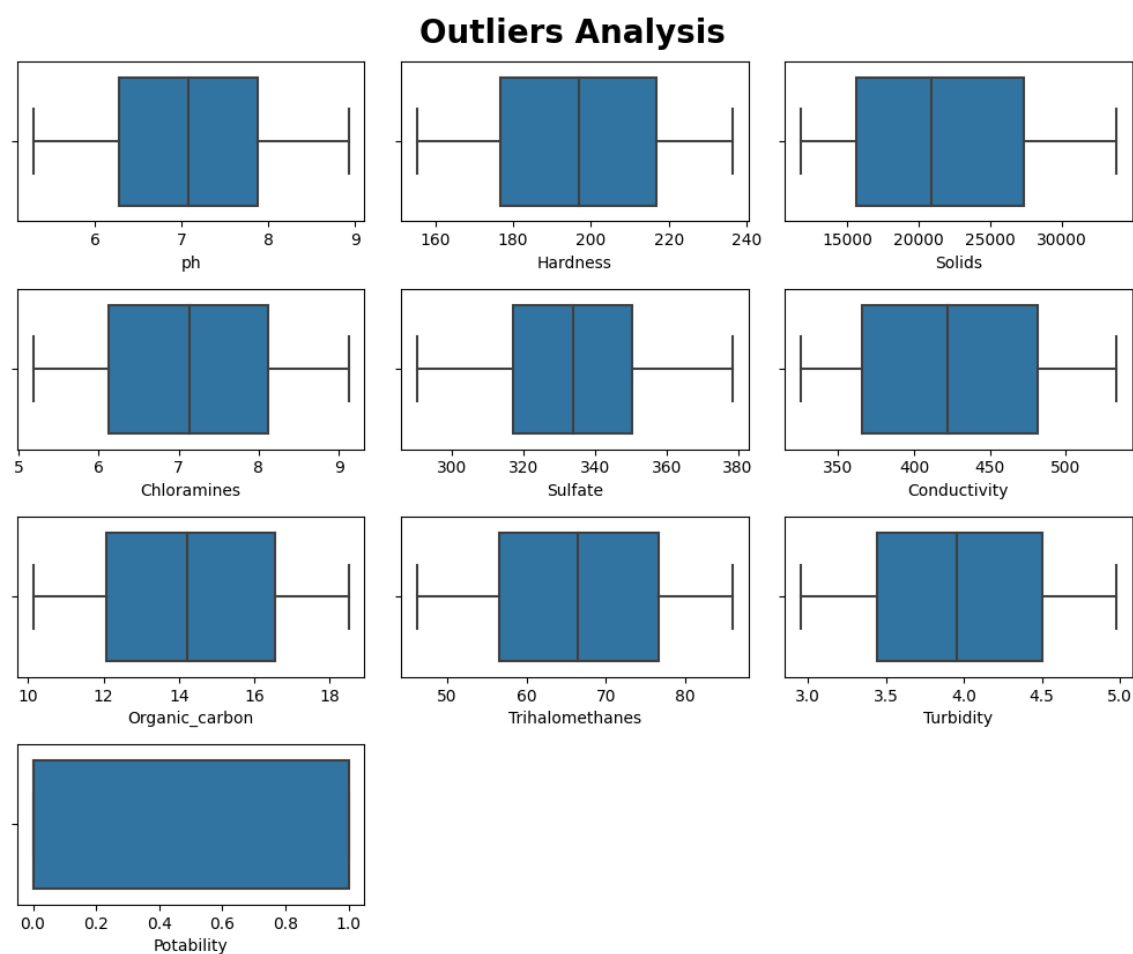
<AxesSubplot: >



Handled Outliers

In [57]:

```
plt.figure(figsize=(10,10))
plt.suptitle("Outliers Analysis", fontsize=20, fontweight="bold")
for i in range(0,len(df.columns)):
    plt.subplot(5,3,i+1)
    sns.boxplot(x=df[df.columns[i]])
    plt.xlabel(df.columns[i])
plt.tight_layout()
```



Feature Selection

FILTER METHOD


1. Pearson's correlation

In [58]:

```
plt.figure(figsize=(15,15))  
sns.heatmap(data=df.corr(), annot=True)
```

Out[58]:

<AxesSubplot: >



fisher's score

In [59]:

```
from skfeature.function.similarity_based import fisher_score
```

In [60]:


```
x = df.drop("Potability", axis=1)  
y = df["Potability"]
```

In [61]:

```
fisher_rank = fisher_score.fisher_score(x.to_numpy(),y)  
s1 = pd.Series(fisher_rank, index= x.columns)  
s1.sort_values().plot(kind="barh")
```

Out[61]:

<AxesSubplot: >



Variance threshold method

In [62]:

```
from sklearn.feature_selection import VarianceThreshold
```

In [63]:

```
var_th = VarianceThreshold(threshold = 0.3)
var_th.fit_transform(df)
arr = var_th.get_support()
np.where(arr == False)
df.columns[np.where(arr == False)]
```

Out[63]:

```
Index(['Potability'], dtype='object')
```

variance inflationn factor

In [64]:

```
from statsmodels.stats.outliers_influence import variance_inflation_factor
```

In [65]:

```
x.values
```

Out[65]:

```
array([[7.08079450e+00, 2.04890455e+02, 2.07913190e+04, ...,
        1.03797831e+01, 8.59000947e+01, 2.96313538e+00],
       [5.28219449e+00, 1.55223964e+02, 1.86300579e+04, ...,
        1.51800131e+01, 5.63290763e+01, 4.50065627e+00],
       [8.09912419e+00, 2.24236259e+02, 1.99095417e+04, ...,
        1.68686369e+01, 6.64200925e+01, 3.05593375e+00],
       ...,
       [8.92504688e+00, 1.75762646e+02, 3.31555782e+04, ...,
        1.10390697e+01, 6.98454003e+01, 3.29887550e+00],
       [5.28219449e+00, 2.30603758e+02, 1.19838694e+04, ...,
        1.11689462e+01, 7.74882131e+01, 4.70865847e+00],
       [7.87467136e+00, 1.95102299e+02, 1.74041771e+04, ...,
        1.61403676e+01, 7.86984463e+01, 2.95180283e+00]])
```

In [66]:

```
vif_lst = []
for i in range(x.shape[1]):
    vif = variance_inflation_factor(x.values,i)
    vif_lst.append(vif)
vif_lst
```

Out[66]:


```
[39.60268140646638,
 50.81765606290291,
 9.854114392291905,
 31.12985728335243,
 98.62440998694287,
 36.253227605395594,
 27.448437065991804,
 26.83919183339501,
 34.67773797589433]
```

In [67]:

```
s1 = pd.Series(vif_lst, index= x.columns)
s1.sort_values()
s1.sort_values().plot(kind= "barh")
```

Out[67]:

<AxesSubplot: >



Information Gain

In [68]:


```
from sklearn.feature_selection import mutual_info_regression
```


In [69]:

```
a1=mutual_info_regression(x,y)
s2 = pd.Series(a1, index= x.columns)
s2.sort_values().plot(kind="barh")
```

Out[69]:

<AxesSubplot: >



sulphate and Trihalomethanes will be dropped as these two features don't contribute in our target column

In [70]:

```
df_backup = df
df_backup.head(2)
```

Out[70]:

	ph	Hardness	Solids	Chloramines	Sulfate	Conductivity	Organic_carbon
0	7.080795	204.890455	20791.318981	7.300212	368.516441	533.297241	10.37978
1	5.282194	155.223964	18630.057858	6.635246	333.775777	533.297241	15.1800

In [71]:

```
df_new = df.drop(["Sulfate", "Trihalomethanes"], axis=1)
```

In [72]:

```
df_new.head(2)
```

Out[72]:

	ph	Hardness	Solids	Chloramines	Conductivity	Organic_carbon	Turbidity
0	7.080795	204.890455	20791.318981	7.300212	533.297241	10.379783	2.963135
1	5.282194	155.223964	18630.057858	6.635246	533.297241	15.180013	4.500656

WRAPPER METHOD

Forward Elimination

In [73]:

```
from sklearn.feature_selection import SequentialFeatureSelector
from sklearn.neighbors import KNeighborsRegressor
import time
```


In [74]:

```
x1 = df.drop("Potability", axis=1)
y1 = df["Potability"]
```

In [75]:

```
strt_tm = time.time()
knn_model = KNeighborsRegressor()
sfs = SequentialFeatureSelector(knn_model, n_features_to_select=5, direction="forward", cv=5)
sfs.fit(x1,y1)
```

Out[75]:



In [76]:

```
end_tm = time.time()
total = end_tm - strt_tm
print("total time taken: ", total)
a2 = sfs.get_support()
s3 = pd.Series(a2, x1.columns)
s3
```

total time taken: 4.894103765487671

Out[76]:

```
ph                False
Hardness          True
Solids            True
Chloramines       True
Sulfate           False
Conductivity      False
Organic_carbon    True
Trihalomethanes   False
Turbidity         True
dtype: bool
```

Backward Elimination

In [77]:

```
strt_tm = time.time()
knn_model = KNeighborsRegressor()
sfs = SequentialFeatureSelector(knn_model, n_features_to_select=5, direction="backward",
sfs.fit(x1,y1)
end_tm = time.time()
total = end_tm - strt_tm
print("total time taken: ", total)
a2 = sfs.get_support()
s3 = pd.Series(a2, x1.columns)
s3
```

total time taken: 0.4389803409576416

Out[77]:

```
ph                False
Hardness          False
Solids            False
Chloramines       True
Sulfate           True
Conductivity      True
Organic_carbon    False
Trihalomethanes   True
Turbidity         True
dtype: bool
```

In [78]:

```
df.drop("Sulfate",axis=1, inplace=True)
```

In [79]:

```
df.head()
```

Out[79]:

	ph	Hardness	Solids	Chloramines	Conductivity	Organic_carbon	Trihalomethanes
0	7.080795	204.890455	20791.318981	7.300212	533.297241	10.379783	85
1	5.282194	155.223964	18630.057858	6.635246	533.297241	15.180013	56
2	8.099124	224.236259	19909.541732	9.122578	418.606213	16.868637	66
3	8.316766	214.373394	22018.417441	8.059332	363.266516	18.436524	85
4	8.925047	181.101509	17978.986339	6.546600	398.410813	11.558279	46

Model Evaluation

with df >> Sulfate column in dropped

In [94]:

```
df.columns
```

Out[94]:

```
Index(['ph', 'Hardness', 'Solids', 'Chloramines', 'Conductivity',
      'Organic_carbon', 'Trihalomethanes', 'Turbidity', 'Potability'],
      dtype='object')
```

In [82]:

```
x = df.drop("Potability",axis=1)
y = df["Potability"]
```

In [83]:

```
x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.2,random_state=42)
```

In [84]:

```
x_train.shape , x_test.shape
```

Out[84]:

```
((2620, 8), (656, 8))
```

In [85]:

```
y_train.shape , y_test.shape
```

Out[85]:

```
((2620,), (656,))
```

Logistic Regression

In [86]:

```
lg_model = LogisticRegression()
lg_model.fit(x_train,y_train)
```

Out[86]:

```
▼ LogisticRegression
LogisticRegression()
```

testing

In [87]:

```
y_pred = lg_model.predict(x_test)
acc_score = accuracy_score(y_test,y_pred)
print("Test_LG_Acc: ",acc_score)

con_matrix = confusion_matrix(y_test,y_pred)
print("Confusion_Matrix:\n " ,con_matrix)

clf_report = classification_report(y_test,y_pred)
print("Classification_report: \n" , clf_report )
```

```
Test_LG_Acc: 0.6280487804878049
```

```
Confusion_Matrix:
```

```
[[412  0]
```

```
[244  0]]
```

```
Classification_report:
```

	precision	recall	f1-score	support
0	0.63	1.00	0.77	412
1	0.00	0.00	0.00	244
accuracy			0.63	656
macro avg	0.31	0.50	0.39	656
weighted avg	0.39	0.63	0.48	656

Training

In [88]:

```
y_pred_train = lg_model.predict(x_train)
acc_score = accuracy_score(y_train,y_pred_train)
print("Train_LG_Acc: ",acc_score)

con_matrix = confusion_matrix(y_train,y_pred_train)
print("Confusion_Matrix:\n " ,con_matrix)

clf_report = classification_report(y_train,y_pred_train)
print("Classification_report: \n" , clf_report )
```

Train_LG_Acc: 0.6053435114503817

Confusion_Matrix:

```
[[1586  0]
```

```
[1034  0]]
```

Classification_report:

	precision	recall	f1-score	support
0	0.61	1.00	0.75	1586
1	0.00	0.00	0.00	1034
accuracy			0.61	2620
macro avg	0.30	0.50	0.38	2620
weighted avg	0.37	0.61	0.46	2620

Random Forest

In [91]:

```
rf_model = RandomForestClassifier()
rf_model.fit(x_train,y_train)
```

Out[91]:

```
▼ RandomForestClassifier
RandomForestClassifier()
```

In [92]:

```

y_pred_test = rf_model.predict(x_test)
conf_matrix = confusion_matrix(y_test,y_pred_test)
print("Confusion Matrix: \n",conf_matrix)

Acc = accuracy_score(y_test,y_pred_test)
print("Accuracy: ",Acc)

Clss_report = classification_report(y_test,y_pred_test)
print("Classification Report: \n", Clss_report)

```

Confusion Matrix:

```

[[357  55]
 [180  64]]

```

Accuracy: 0.6417682926829268

Classification Report:

	precision	recall	f1-score	support
0	0.66	0.87	0.75	412
1	0.54	0.26	0.35	244
accuracy			0.64	656
macro avg	0.60	0.56	0.55	656
weighted avg	0.62	0.64	0.60	656

In [93]:

```

y_pred_train = rf_model.predict(x_train)
conf_matrix = confusion_matrix(y_train,y_pred_train)
print("Confusion Matrix: \n",conf_matrix)

Acc = accuracy_score(y_train,y_pred_train)
print("Accuracy: ",Acc)

Clss_report = classification_report(y_train,y_pred_train)
print("Classification Report: \n", Clss_report)

```

Confusion Matrix:

```

[[1586   0]
 [   0 1034]]

```

Accuracy: 1.0

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	1586
1	1.00	1.00	1.00	1034
accuracy			1.00	2620
macro avg	1.00	1.00	1.00	2620
weighted avg	1.00	1.00	1.00	2620

with df_new >> after feature selection sulfate, Trihalomethanes was dropped

In [97]:

```
df_new.columns
```

Out[97]:

```
Index(['ph', 'Hardness', 'Solids', 'Chloramines', 'Conductivity',  
      'Organic_carbon', 'Turbidity', 'Potability'],  
      dtype='object')
```

In [98]:

```
x1 = df_new.drop("Potability",axis=1)  
y1 = df_new["Potability"]
```

In [99]:

```
x_train,x_test,y_train,y_test = train_test_split(x1,y1,test_size=0.2,random_state=42)
```

Logistic Regression

In [100]:

```
lg_model = LogisticRegression()  
lg_model.fit(x_train,y_train)
```

Out[100]:

```
▼ LogisticRegression  
LogisticRegression()
```


In [101]:

```

y_pred = lg_model.predict(x_test)
acc_score = accuracy_score(y_test,y_pred)
print("Test_LG_Acc: ",acc_score)

con_matrix = confusion_matrix(y_test,y_pred)
print("Confusion_Matrix:\n ",con_matrix)

clf_report = classification_report(y_test,y_pred)
print("Classification_report: \n" , clf_report )

```

Test_LG_Acc: 0.6280487804878049

Confusion_Matrix:

[[412 0]

[244 0]]

Classification_report:

	precision	recall	f1-score	support
0	0.63	1.00	0.77	412
1	0.00	0.00	0.00	244
accuracy			0.63	656
macro avg	0.31	0.50	0.39	656
weighted avg	0.39	0.63	0.48	656

In [102]:

```

y_pred_train = lg_model.predict(x_train)
acc_score = accuracy_score(y_train,y_pred_train)
print("Train_LG_Acc: ",acc_score)

con_matrix = confusion_matrix(y_train,y_pred_train)
print("Confusion_Matrix:\n ",con_matrix)

clf_report = classification_report(y_train,y_pred_train)
print("Classification_report: \n" , clf_report )

```

Train_LG_Acc: 0.6053435114503817

Confusion_Matrix:

[[1586 0]

[1034 0]]

Classification_report:

	precision	recall	f1-score	support
0	0.61	1.00	0.75	1586
1	0.00	0.00	0.00	1034
accuracy			0.61	2620
macro avg	0.30	0.50	0.38	2620
weighted avg	0.37	0.61	0.46	2620

df_backup >> all the features are present

In [103]:

```
df_backup.columns
```

Out[103]:

```
Index(['ph', 'Hardness', 'Solids', 'Chloramines', 'Conductivity',  
      'Organic_carbon', 'Trihalomethanes', 'Turbidity', 'Potability'],  
      dtype='object')
```

In [104]:

```
x2 = df_backup.drop("Potability",axis=1)  
y2 = df_backup["Potability"]
```

In [109]:

```
x2_train,x2_test,y2_train,y2_test = train_test_split(x2,y2,test_size=0.2,random_state=42)
```

In [110]:

```
lg_model = LogisticRegression()  
lg_model.fit(x2_train,y2_train)
```

Out[110]:

```
▼ LogisticRegression  
LogisticRegression()
```

In [112]:

```

y2_pred = lg_model.predict(x2_test)
acc_score = accuracy_score(y2_test,y2_pred)
print("Test_LG_Acc: ",acc_score)

con_matrix = confusion_matrix(y2_test,y2_pred)
print("Confusion_Matrix:\n ",con_matrix)

clf_report = classification_report(y2_test,y2_pred)
print("Classification_report: \n" , clf_report )

```

Test_LG_Acc: 0.6280487804878049

Confusion_Matrix:

[[412 0]

[244 0]]

Classification_report:

	precision	recall	f1-score	support
0	0.63	1.00	0.77	412
1	0.00	0.00	0.00	244
accuracy			0.63	656
macro avg	0.31	0.50	0.39	656
weighted avg	0.39	0.63	0.48	656

In [113]:

```

y2_pred_train = lg_model.predict(x2_train)
acc_score = accuracy_score(y2_train,y2_pred_train)
print("Train_LG_Acc: ",acc_score)

con_matrix = confusion_matrix(y2_train,y2_pred_train)
print("Confusion_Matrix:\n ",con_matrix)

clf_report = classification_report(y2_train,y2_pred_train)
print("Classification_report: \n" , clf_report )

```

Train_LG_Acc: 0.6053435114503817

Confusion_Matrix:

[[1586 0]

[1034 0]]

Classification_report:

	precision	recall	f1-score	support
0	0.61	1.00	0.75	1586
1	0.00	0.00	0.00	1034
accuracy			0.61	2620
macro avg	0.30	0.50	0.38	2620
weighted avg	0.37	0.61	0.46	2620

Random Forest

In [119]:

```
rf_model = RandomForestClassifier()
rf_model.fit(x2_train,y2_train)
```

Out[119]:

```
▼ RandomForestClassifier
RandomForestClassifier()
```

In [120]:

```
y2_pred_test = rf_model.predict(x2_test)
conf_matrix = confusion_matrix(y2_test,y2_pred_test)
print("Confusion Matrix: \n",conf_matrix)

Acc = accuracy_score(y2_test,y2_pred_test)
print("Accuracy: ",Acc)

Clss_report = classification_report(y2_test,y2_pred_test)
print("Classification Report: \n", Clss_report)
```

Confusion Matrix:

```
[[350  62]
 [184  60]]
```

Accuracy: 0.625

Classification Report:

	precision	recall	f1-score	support
0	0.66	0.85	0.74	412
1	0.49	0.25	0.33	244
accuracy			0.62	656
macro avg	0.57	0.55	0.53	656
weighted avg	0.59	0.62	0.59	656

In [121]:

```

y2_pred_train = rf_model.predict(x2_train)
conf_matrix = confusion_matrix(y2_train,y2_pred_train)
print("Confusion Matrix: \n",conf_matrix)

Acc = accuracy_score(y2_train,y2_pred_train)
print("Accuracy: ",Acc)

Cls_report = classification_report(y2_train,y2_pred_train)
print("Classification Report: \n", Cls_report)

```

Confusion Matrix:

```

[[1586   0]
 [   0 1034]]

```

Accuracy: 1.0

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	1586
1	1.00	1.00	1.00	1034
accuracy			1.00	2620
macro avg	1.00	1.00	1.00	2620
weighted avg	1.00	1.00	1.00	2620

hyperparameter tuning


In [124]:

```

hyperparameters = {"n_estimators": np.arange(10,100),
                    "criterion": ["gini","entropy"],
                    "max_depth": np.arange(3,8),
                    "min_samples_split": np.arange(2,20),
                    "min_samples_leaf": np.arange(2,15),
                    "random_state": [11],
                    "oob_score": [True],
                    "max_features": ["auto"]}
rscv_rf_model = RandomizedSearchCV(rf_model,hyperparameters,cv=7)
rscv_rf_model.fit(x2_train,y2_train)

```

Out[124]:



In [125]:

```

y2_pred_test = rscv_rf_model.predict(x2_test)
conf_matrix = confusion_matrix(y2_test,y2_pred_test)
print("Confusion Matrix: \n",conf_matrix)

Test_RF_Acc = accuracy_score(y2_test,y2_pred_test)
print("Accuracy: ",Test_RF_Acc)

Cls_report = classification_report(y2_test,y2_pred_test)
print("Classification Report: \n", Cls_report)

```

Confusion Matrix:

```

[[401  11]
 [224  20]]

```

Accuracy: 0.6417682926829268

Classification Report:

	precision	recall	f1-score	support
0	0.64	0.97	0.77	412
1	0.65	0.08	0.15	244
accuracy			0.64	656
macro avg	0.64	0.53	0.46	656
weighted avg	0.64	0.64	0.54	656

In [126]:

```

y2_pred_train = rscv_rf_model.predict(x2_train)
conf_matrix = confusion_matrix(y2_train,y2_pred_train)
print("Confusion Matrix: \n",conf_matrix)

Train_RF_Acc = accuracy_score(y2_train,y2_pred_train)
print("Accuracy: ",Train_RF_Acc)

Cls_report = classification_report(y2_train,y2_pred_train)
print("Classification Report: \n", Cls_report)

```

Confusion Matrix:

```

[[1575  11]
 [ 847 187]]

```

Accuracy: 0.6725190839694657

Classification Report:

	precision	recall	f1-score	support
0	0.65	0.99	0.79	1586
1	0.94	0.18	0.30	1034
accuracy			0.67	2620
macro avg	0.80	0.59	0.54	2620
weighted avg	0.77	0.67	0.60	2620

**finally got accuracy with original data tessting = 64 and
trAining=67**

In []: